

ECE Selected Topics in Machine Learning – Assignment 4

For CIFAR-10 we were told to achieve a performance similar to the state of the art, but I found papers that claimed accuracies of above 98%. I quickly came to realize this was not going to be possible for me for this assignment. After playing with various features, implementing datagen/ImageDataGenerator, normalization techniques as well as LearningRateScheduler I was able to achieve accuracies in the mid 80s range. I was able to reuse the basic structure of the mnist program which I started with, but I had to add a lot more capabilities in order to make it work for CIFAR with a reasonable accuracy. I was able to mostly use the same code for CIFAR-100 from CIFAR-10 with minor modifications, as only 80% was needed for top 5 accuracy.

The datagen/ImageDataGenerator helped perform data augmentation, which was helpful in increasing the accuracy of the model as the diversity of the data used while training the model was increased. Normalization and changing the learning rates were also new methods I learned improved the accuracy of the model. I did not play with the optimizer as Adam suited these purposes, but for further improvement, this is an area to look more into. I had trouble with this assignment with overfitting, as the validation accuracy repeatedly came out low; however, when I decreased the complexity of the model, this resulted in a decrease in test case accuracy, so I chose to leave the model with low validation accuracy. I did the shuffle the data to make sure it was not an ordering issue.

Here is the Program for CIFAR-10:

```
import tensorflow as tf
import numpy as np
from sklearn.model_selection import train_test_split

LAMBDA = .0001
NUM_CLASSES = 10
BATCH_SIZE = 64

def lr_schedule(epoch):
    lr = 0.001
    if epoch > 75:
        lr = 0.0005
    elif epoch > 100:
        lr = 0.00025
    elif epoch > 150:
        lr = 0.000056
    return lr

class Data(object):
    def __init__(self):

        (self.x_train, self.y_train), (self.x_test, self.y_test) = tf.keras.datasets.cifar10.load_data()

        rand_index = np.arange(len(self.x_train))
        np.random.shuffle(rand_index)
        rand_index2 = np.arange(len(self.x_test))
        np.random.shuffle(rand_index2)
        self.x_train = self.x_train[rand_index]
        self.y_train = self.y_train[rand_index]
        self.x_test = self.x_test[rand_index2]
        self.y_test = self.y_test[rand_index2]
```

```

self.x_train, self.x_ver, self.y_train, self.y_ver = train_test_split(self.x_train, self.y_train, test_size=5000)

self.x_train = self.x_train.reshape(self.x_train.shape[0], 32, 32, 3).astype('float32')/255.0
self.x_ver = self.x_ver.reshape(self.x_ver.shape[0], 32, 32, 3).astype('float32')/255.0
self.x_test = self.x_test.reshape(self.x_test.shape[0], 32, 32, 3).astype('float32')/255.0

mean = np.mean(self.x_train,axis=(0,1,2,3))
std = np.std(self.x_train,axis=(0,1,2,3))
self.x_train = (self.x_train-mean)/(std+1e-7)
self.x_test = (self.x_test-mean)/(std+1e-7)

self.y_train = tf.keras.utils.to_categorical(self.y_train, NUM_CLASSES)
self.y_test = tf.keras.utils.to_categorical(self.y_test, NUM_CLASSES)
self.y_ver = tf.keras.utils.to_categorical(self.y_ver, NUM_CLASSES)

class Model(tf.Module):
    def __init__(self):
        self.model = tf.keras.Sequential()
        self.model.add(tf.keras.layers.Conv2D(32, (3,3), padding='same', kernel_regularizer=tf.keras.regularizers.l2(LAMBDA),
input_shape=data.x_train.shape[1:]))

        def activation(self):
            self.model.add(tf.keras.layers.Activation('elu'))

        def batch_norm(self):
            self.model.add(tf.keras.layers.BatchNormalization())

        def conv2D_layer(self):
            self.model.add(tf.keras.layers.Conv2D(128, (3,3), padding='same', kernel_regularizer=tf.keras.regularizers.l2(LAMBDA)))

        def maxPooling2D(self):
            self.model.add(tf.keras.layers.MaxPooling2D(pool_size=(2,2)))

if __name__ == "__main__":
    data = Data()
    model = Model()

    model.activation()
    model.batch_norm()

    model.conv2D_layer()
    model.activation()
    model.batch_norm()

    model.maxPooling2D()
    model.model.add(tf.keras.layers.Dropout(0.4))

    model.conv2D_layer()
    model.activation()
    model.batch_norm()

    model.conv2D_layer()
    model.activation()
    model.batch_norm()

    model.maxPooling2D()

    model.model.add(tf.keras.layers.Dropout(0.4))
    model.model.add(tf.keras.layers.Flatten())
    model.model.add(tf.keras.layers.Dense(NUM_CLASSES, activation='softmax'))
    model.model.summary()

```

```

datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
)

datagen.fit(data.x_train)

model.model.compile(loss='categorical_crossentropy',
    optimizer=tf.keras.optimizers.Adam(),
    metrics=['accuracy', 'top_k_categorical_accuracy'])

model.model.fit_generator(datagen.flow(data.x_train, data.y_train, batch_size=BATCH_SIZE),
    use_multiprocessing=True,
    steps_per_epoch=data.x_train.shape[0] // BATCH_SIZE, epochs=100,
    verbose=1, validation_data=(data.x_ver, data.y_ver),
    callbacks=[tf.keras.callbacks.LearningRateScheduler(lr_schedule)], workers = 4)

final_result = model.model.evaluate(data.x_test, data.y_test, batch_size=128, verbose=1)
print("Test Loss:", final_result[0])
print("Test Accuracy:", final_result[1])
print("Test Top-5 Accuracy:", final_result[2])

```

Output for 150 epochs:

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 128)	36992
activation_1 (Activation)	(None, 32, 32, 128)	0
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 128)	512
max_pooling2d (MaxPooling2D)	(None, 16, 16, 128)	0
dropout (Dropout)	(None, 16, 16, 128)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	147584
activation_2 (Activation)	(None, 16, 16, 128)	0
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 128)	512
conv2d_3 (Conv2D)	(None, 16, 16, 128)	147584
activation_3 (Activation)	(None, 16, 16, 128)	0
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)	0
dropout_1 (Dropout)	(None, 8, 8, 128)	0
flatten (Flatten)	(None, 8192)	0

dense (Dense) (None, 10) 81930

=====

Total params: 416,650
Trainable params: 415,818
Non-trainable params: 832

Epoch 98/100

702/703 [=====>.] - ETA: 0s - loss: 0.4709 - acc: 0.8784 - top_k_categorical_accuracy: 0.9964
703/703 [=====] - 11s 15ms/step - loss: 0.4708 - acc: 0.8784 - top_k_categorical_accuracy: 0.9964 - val_loss: 2.1340
- val_acc: 0.3979 - val_top_k_categorical_accuracy: 0.8145

Epoch 99/100

702/703 [=====>.] - ETA: 0s - loss: 0.4710 - acc: 0.8770 - top_k_categorical_accuracy: 0.9970
703/703 [=====] - 11s 15ms/step - loss: 0.4711 - acc: 0.8770 - top_k_categorical_accuracy: 0.9970 - val_loss: 2.1725
- val_acc: 0.4290 - val_top_k_categorical_accuracy: 0.8064

Epoch 100/100

702/703 [=====>.] - ETA: 0s - loss: 0.4606 - acc: 0.8804 - top_k_categorical_accuracy: 0.9969
703/703 [=====] - 11s 15ms/step - loss: 0.4606 - acc: 0.8804 - top_k_categorical_accuracy: 0.9970 - val_loss: 2.0650
- val_acc: 0.4316 - val_top_k_categorical_accuracy: 0.8200

10000/10000 [=====] - 1s 71us/sample - loss: 0.5561 - acc: 0.8664 - top_k_categorical_accuracy: 0.9943

Test Loss: 0.5560953780651092

Test Accuracy: 0.8664

Test Top-5 Accuracy: 0.9943

For the CIFAR-100, we were told to achieve a performance of top-5 accuracy of 80%.

Here is the program for CIFAR-100:

```
import tensorflow as tf
import numpy as np
from sklearn.model_selection import train_test_split

LAMBDA = .0001
NUM_CLASSES = 100
BATCH_SIZE = 64

class Data(object):
    def __init__(self):

        (self.x_train, self.y_train), (self.x_test, self.y_test) = tf.keras.datasets.cifar100.load_data()

        rand_index = np.arange(len(self.x_train))
        np.random.shuffle(rand_index)
        rand_index2 = np.arange(len(self.x_test))
        np.random.shuffle(rand_index2)
        self.x_train = self.x_train[rand_index]
        self.y_train = self.y_train[rand_index]
        self.x_test = self.x_test[rand_index2]
        self.y_test = self.y_test[rand_index2]

        self.x_train, self.x_ver, self.y_train, self.y_ver = train_test_split(self.x_train, self.y_train, test_size=5000)

        self.x_train = self.x_train.reshape(self.x_train.shape[0], 32, 32, 3).astype('float32')/255.0
        self.x_ver = self.x_ver.reshape(self.x_ver.shape[0], 32, 32, 3).astype('float32')/255.0
        self.x_test = self.x_test.reshape(self.x_test.shape[0], 32, 32, 3).astype('float32')/255.0

        mean = np.mean(self.x_train,axis=(0,1,2,3))
        std = np.std(self.x_train,axis=(0,1,2,3))
        self.x_train = (self.x_train-mean)/(std+1e-7)
        self.x_test = (self.x_test-mean)/(std+1e-7)

        self.y_train = tf.keras.utils.to_categorical(self.y_train, NUM_CLASSES)
        self.y_test = tf.keras.utils.to_categorical(self.y_test, NUM_CLASSES)
        self.y_ver = tf.keras.utils.to_categorical(self.y_ver, NUM_CLASSES)
```

```

class Model(tf.Module):
    def __init__(self):
        self.model = tf.keras.Sequential()
        self.model.add(tf.keras.layers.Conv2D(32, (3,3), padding='same', kernel_regularizer=tf.keras.regularizers.l2(LAMBDA),
input_shape=data.x_train.shape[1:]))

    def activation(self):
        self.model.add(tf.keras.layers.Activation('elu'))

    def batch_norm(self):
        self.model.add(tf.keras.layers.BatchNormalization())

    def conv2D_layer(self):
        self.model.add(tf.keras.layers.Conv2D(128, (3,3), padding='same', kernel_regularizer=tf.keras.regularizers.l2(LAMBDA)))

    def maxPooling2D(self):
        self.model.add(tf.keras.layers.MaxPooling2D(pool_size=(2,2)))

if __name__ == "__main__":
    data = Data()
    model = Model()

    model.activation()
    model.batch_norm()

    model.conv2D_layer()
    model.activation()
    model.batch_norm()

    model.maxPooling2D()
    model.model.add(tf.keras.layers.Dropout(0.4))

    model.conv2D_layer()
    model.activation()
    model.batch_norm()

    model.conv2D_layer()
    model.activation()
    model.batch_norm()

    model.maxPooling2D()

    model.model.add(tf.keras.layers.Dropout(0.4))
    model.model.add(tf.keras.layers.Flatten())
    model.model.add(tf.keras.layers.Dense(NUM_CLASSES, activation='softmax'))
    model.model.summary()

    datagen = tf.keras.preprocessing.image.ImageDataGenerator(
        rotation_range=15,
        width_shift_range=0.1,
        height_shift_range=0.1,
        horizontal_flip=True,
    )

    datagen.fit(data.x_train)

    model.model.compile(loss='categorical_crossentropy',
optimizer=tf.keras.optimizers.Adam(),
metrics=['accuracy', 'top_k_categorical_accuracy'])

```

```

model.model.fit_generator(datagen.flow(data.x_train, data.y_train, batch_size=BATCH_SIZE),\
    steps_per_epoch=data.x_train.shape[0] // BATCH_SIZE, epochs=50,\
    verbose=1,validation_data=(data.x_ver, data.y_ver))

final_result = model.model.evaluate(data.x_test, data.y_test, batch_size=128, verbose=1)
print("Test Loss:", final_result[0])
print("Test Accuracy:", final_result[1])
print("Test Top-5 Accuracy:", final_result[2])

```

Output for 50 epochs:

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 128)	36992
activation_1 (Activation)	(None, 32, 32, 128)	0
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 128)	512
max_pooling2d (MaxPooling2D)	(None, 16, 16, 128)	0
dropout (Dropout)	(None, 16, 16, 128)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	147584
activation_2 (Activation)	(None, 16, 16, 128)	0
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 128)	512
conv2d_3 (Conv2D)	(None, 16, 16, 128)	147584
activation_3 (Activation)	(None, 16, 16, 128)	0
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)	0
dropout_1 (Dropout)	(None, 8, 8, 128)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 100)	819300
=====		
Total params: 1,154,020		
Trainable params: 1,153,188		
Non-trainable params: 832		

Epoch 45/50

703/703 [=====] - 22s 31ms/step - loss: 1.4527 - acc: 0.6435 - top_k_categorical_accuracy: 0.8979 - val_loss: 4.5644 - val_acc: 0.1474 - val_top_k_categorical_accuracy: 0.3726

Epoch 46/50

703/703 [=====] - 22s 31ms/step - loss: 1.4448 - acc: 0.6410 - top_k_categorical_accuracy: 0.8982 - val_loss: 4.6112 - val_acc: 0.1428 - val_top_k_categorical_accuracy: 0.3560

Epoch 47/50

703/703 [=====] - 22s 31ms/step - loss: 1.4377 - acc: 0.6486 - top_k_categorical_accuracy: 0.8980 - val_loss: 4.7397 - val_acc: 0.1446 - val_top_k_categorical_accuracy: 0.3434

Epoch 48/50

703/703 [=====] - 22s 31ms/step - loss: 1.4294 - acc: 0.6478 - top_k_categorical_accuracy: 0.9005 - val_loss: 4.3527 - val_acc: 0.1674 - val_top_k_categorical_accuracy: 0.3814

Epoch 49/50

703/703 [=====] - 22s 31ms/step - loss: 1.4373 - acc: 0.6490 - top_k_categorical_accuracy: 0.8974 - val_loss: 4.2862
- val_acc: 0.1546 - val_top_k_categorical_accuracy: 0.3858
Epoch 50/50
703/703 [=====] - 22s 31ms/step - loss: 1.4249 - acc: 0.6511 - top_k_categorical_accuracy: 0.8991 - val_loss: 4.8304
- val_acc: 0.1320 - val_top_k_categorical_accuracy: 0.3510
10000/10000 [=====] - 1s 74us/sample - loss: 1.8795 - acc: 0.5669 - top_k_categorical_accuracy: 0.8398
Test Loss: 1.8794701976776123
Test Accuracy: 0.5669
Test Top-5 Accuracy: 0.8398

Sources that helped me! (along with TensorFlow Docs)

<https://www.cs.toronto.edu/~kriz/cifar.html>

<https://www.codecademy.com/articles/normalization>

<https://towardsdatascience.com/cifar-10-image-classification-in-tensorflow-5b501f7dc77c>

<https://machinelearningmastery.com/using-learning-rate-schedules-deep-learning-models-python-keras/>

<https://appliedmachinelearning.blog/author/abhijeetchar/>